




Note: Each cell is an estimated skill parameter Theta: the skill of a worker skill set (column) at a task (row), in the same units as the paper's skill-parameter tables. Unconstrained estimates are the baseline; constrained estimates add the requirement that skill sets can be ranked by absolute advantage, i.e. that one ordering exists in which each skill set is at least as skilled as the previous one at every task. Columns are ordered least- to most-skilled under that ranking, so under the constraint values never fall from left to right; skill sets are numbered as in the paper and the ordering differs by county. Printed numbers are the estimates themselves; colour is rescaled within each county and saturates at that county's 90th percentile of $|\text{Theta}|$, so that a single extreme cell does not wash out the rest. Colours are therefore comparable within a county but not across counties. Each heading reports the Frobenius distance between that county's unconstrained and constrained matrices (the square root of the summed squared cell-by-cell changes), with that distance as a percentage of the unconstrained matrix's own Frobenius norm.